



Applied Machine Learning

Introduction


Kevyn Collins-Thompson

**Associate Professor of Information & Computer Science
University of Michigan**

What is Machine Learning (ML)?

- **The study of computer programs (algorithms) that can learn by example**
- **ML algorithms can generalize from existing examples of a task**
 - e.g. after seeing a *training set of labeled images*, an *image classifier can figure out how to apply labels accurately to new, previously unseen images*

Speech Recognition



Machine Learning models can learn by example

- **Algorithms learn rules from labelled examples**
- A set of labelled examples used for learning is called training data.
- The learned rules should also be able to generalize to correctly recognize or predict new examples not in the training set.

Audio signal



Output text

How do I
get to Ann
Arbor?




Hello!



Please order
me a pizza.

Machine Learning models learn from experience


- **Labeled examples
(Email spam detection)**
- **User feedback
(Clicks on a search page)**
- **Surrounding environment
(self-driving cars)**



Machine Learning brings together statistics, computer science, and more..

- **Statistical methods**
 - *Infer conclusions from data*
 - *Estimate reliability of predictions*
- **Computer science**
 - *Large-scale computing architectures*
 - *Algorithms for capturing, manipulating, indexing, combining, retrieving and performing predictions on data*
 - *Software pipelines that manage the complexity of multiple subtasks*
- **Economics, biology, psychology**
 - *How can an individual or system efficiently improve their performance in a given environment?*
 - *What is learning and how can it be optimized?*

Machine Learning for fraud detection and credit scoring



Web search: query spell-checking, result ranking, content classification and selection, advertising placement

vacations in michigan

All Maps Shopping News Images More ▾ Search tools

Michigan / Destinations

Detroit Cars, Motown & Detroit Institute of Arts		Grand Rapids Parks, gardens, history, beer, sports		Petoskey Lighthouses, marinas, fishing, parks, beaches	
Mackinac Island Lighthouses, caves, lakes		Mackinaw City Lighthouses, zip-lining, history, lakes, parks		Port Huron Shopping, Thomas Edison, beaches, parks, lighthouses	
Traverse City Beaches, wineries, vineyards, shopping, autumn leaf colors		Ann Arbor Parks, shopping, museums, sports, gardens		Holland Beaches, shopping, churches, art, concerts	


Looking for a Weekend Getaway - Visit The Henry Ford Museum
Ad www.thehenryford.org/ ▾
Experience more, save more. Great values offering up to 30% in savings
📍 20900 Oakwood Blvd, Dearborn, Michigan - Closed now - Hours ▾

Beachtowns, Vacations and Packages Near the ... - Pure Michigan
www.michigan.org/hot-spots/beachtowns/ ▾
Getaway to the Michigan beaches and sand dunes of Grand Haven, Holland, South Haven, St. Joseph, Muskegon, Silver Lake Sand Dunes and Saugatuck.



Map data ©2016 Google, INEGI

Machine Learning for Speech Recognition




Machine Learning algorithms are at the heart of the information economy

- Finance: fraud detection, credit scoring
- Web search
- Speech recognition
- eCommerce: Product recommendations
- Email spam filtering
- Health applications: drug design and discovery
- Education: Automated essay scoring

What is Applied Machine Learning?

- Understand basic ML concepts and workflow
- How to properly apply 'black-box' machine learning components and features
- Learn how to apply machine learning algorithms in Python using the scikit-learn package
- What is not covered in this course:
 - *Underlying theory of statistical machine learning*
 - *Lower-level details of how particular ML components work*
 - *In-depth material on more advanced concepts like deep learning*

Recommended text for this course



Introduction to Machine Learning with Python
A Guide for Data Scientists
By Andreas C. Müller and Sarah Guido

O'Reilly Media



Applied Machine Learning

Key Concepts in Machine Learning

Kevyn Collins-Thompson

Associate Professor of Information & Computer Science
University of Michigan

Key types of Machine Learning problems

Supervised machine learning: Learn to predict **target values** from labelled data.

- Classification (target values are discrete classes)
- Regression (target values are continuous values)

Supervised Learning (classification example)

Training set

X Sample	Y Target Value (Label)		
	x_1	Apple	y_1
	x_2	Lemon	y_2
	x_3	Apple	y_3
	x_4	Orange	y_4

Classifier
 $f : X \rightarrow Y$



At training time, the classifier uses labelled examples to learn rules for recognizing each fruit type.

Future sample




Label: Orange


After training, at prediction time, the trained model is used to predict the fruit type for new instances using the learned rules.

Examples of explicit and implicit label sources

Explicit labels



Implicit labels



Clicking and reading the "Mackinac Island" result can be an implicit label for the search engine to learn that "Mackinac Island" is especially relevant for the query [vacations in michigan] for that specific user.

Key types of Machine Learning problems

Supervised machine learning: Learn to predict target values from labelled data.



- Classification (target values are discrete classes)
- Regression (target values are continuous values)

Unsupervised machine learning: Find structure in *unlabeled data*


- Find groups of similar instances in the data (clustering)
- Finding unusual patterns (outlier detection)

Unsupervised learning: finding useful structure or knowledge in data when no labels are available

- Finding clusters of similar users (clustering)
- Detecting abnormal server access patterns (unsupervised outlier detection)



A Basic Machine Learning Workflow



Representation

Choose:

- A feature representation
- Type of classifier to use

e.g. image pixels, with
k-nearest neighbor classifier

Evaluation

Choose:

- What criterion
distinguishes good vs. bad
classifiers?

e.g. % correct predictions on test set

Optimization

Choose:

- How to search for the
settings/parameters that
give the best classifier
for this evaluation criterion

e.g. try a range of values for "k" parameter
in k-nearest neighbor classifier

Feature Representations

Email

```
To: Chris Brooks
From: Daniel Romero
Subject: Next course offering
Hi Daniel,
Could you please send the outline for the
next course offering? Thanks! -- Chris
```




Feature	Count
to	1
chris	2
brooks	1
from	1
daniel	2
romero	1
the	2
...	

Feature representation

A list of words with their frequency counts

Picture



A matrix of color values (pixels)


Sea Creatures



Feature	Value
DorsalFin	Yes
MainColor	Orange
Stripes	Yes
StripeColor1	White
StripeColor2	Black
Length	4.3 cm

A set of attribute values

Representing a piece of fruit as an array of features (plus label information)



1. Feature representation

Label information
(available in training data only)

Feature representation

	fruit_label	fruit_name	fruit_subtype	mass	width	height	color_score
18	1	apple	cripps_pink	162	7.5	7.1	0.83

2. Learning model

Classifier

Predicted class
(apple)

Represent / Train / Evaluate / Refine Cycle

Representation:

Extract and
select object
features

Represent / Train / Evaluate / Refine Cycle

Representation:


Extract and
select object
features




Train models:

Fit the estimator
to the data

Represent / Train / Evaluate / Refine Cycle



Represent / Train / Evaluate / Refine Cycle





Applied Machine Learning

Python Tools for Machine Learning

Kevyn Collins-Thompson

Associate Professor of Information & Computer Science
University of Michigan

scikit-learn: Python Machine Learning

Library

- **scikit-learn Homepage**
<http://scikit-learn.org/>
- **scikit-learn User Guide**
http://scikit-learn.org/stable/user_guide.html
- **scikit-learn API reference**
<http://scikit-learn.org/stable/modules/classes.html>
- **In Python, we typically import classes and functions we need like this:**

```
from sklearn.model_selection import train_test_split  
from sklearn.tree import DecisionTreeClassifier
```




SciPy Library: Scientific Computing Tools



<http://www.scipy.org/>

- Provides a variety of useful scientific computing tools, including statistical distributions, optimization of functions, linear algebra, and a variety of specialized mathematical functions.
- With scikit-learn, it provides support for *sparse matrices*, a way to store large tables that consist mostly of zeros.
- Example import: `import scipy as sp`

NumPy: Scientific Computing Library




<http://www.numpy.org/>

- Provides fundamental data structures used by scikit-learn, particularly multi-dimensional arrays.
- Typically, data that is input to scikit-learn will be in the form of a NumPy array.
- Example import: `import numpy as np`

Pandas: Data Manipulation and

pandas

$$y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$$



S

<http://pandas.pydata.org/>

- Provides key data structures like DataFrame
- Also, support for reading/writing data in different formats
- Example import: `import pandas as pd`

matplotlib and other plotting libraries



<http://matplotlib.org/>

- We typically use matplotlib's **pyplot** module:
`import matplotlib.pyplot as plt`
- We also sometimes use the **seaborn** visualization library (<http://seaborn.pydata.org/>)
`import seaborn as sn`
- And sometimes the **graphviz** plotting library:
`import graphviz`

Versions of main libraries used in this course

Library name	Minimum version
scikit-learn	0.18.1
scipy	0.19.0
numpy	1.12.1
pandas	0.19.2
matplotlib	2.0.1
seaborn	0.7.1
graphviz	0.7

It's okay if your versions of these don't match ours exactly, as long as the version of scikit-learn and other libraries you're using is the same or greater than listed here.




Applied Machine Learning

An Example Machine Learning Problem

Kevyn Collins-Thompson

Associate Professor of Information & Computer Science
University of Michigan

The Fruit Dataset



	fruit_label	fruit_name	fruit_subtype	mass	width	height	color_score
0	1	apple	granny_smith	192	8.4	7.3	0.55
1	1	apple	granny_smith	180	8.0	6.8	0.59
2	1	apple	granny_smith	176	7.4	7.2	0.60
3	2	mandarin	mandarin	86	6.2	4.7	0.80
4	2	mandarin	mandarin	84	6.0	4.6	0.79
5	2	mandarin	mandarin	80	5.8	4.3	0.77
6	2	mandarin	mandarin	80	5.9	4.3	0.81
7	2	mandarin	mandarin	76	5.8	4.0	0.81
8	1	apple	braeburn	178	7.1	7.8	0.92
9	1	apple	braeburn	172	7.4	7.0	0.89
10	1	apple	braeburn	166	6.9	7.3	0.93
11	1	apple	braeburn	172	7.1	7.6	0.92
12	1	apple	braeburn	154	7.0	7.1	0.88
13	1	apple	golden_delicious	164	7.3	7.7	0.70
14	1	apple	golden_delicious	152	7.6	7.3	0.69
15	1	apple	golden_delicious	156	7.7	7.1	0.69
16	1	apple	golden_delicious	156	7.6	7.5	0.67

`fruit_data_with_colors.txt`

Credit: Original version of the fruit dataset created by Dr. Iain Murray, Univ. of Edinburgh

The input data as a table


Each row corresponds to a single data instance (sample)

The fruit_label column contains the label for each data instance (sample)


	fruit_label	fruit_name	fruit_subtype	mass	width	height	color_score
0	1	apple	granny_smith	192	8.4	7.3	0.55
1	1	apple	granny_smith	180	8.0	6.9	0.59
2	1	apple	granny_smith	176	7.4	7.2	0.60
3	2	mandarin	mandarin	86	6.2	4.7	0.80
4	2	mandarin	mandarin	84	6.0	4.6	0.79
5	2	mandarin	mandarin	80	5.8	4.3	0.77
6	2	mandarin	mandarin	80	5.9	4.3	0.81
7	2	mandarin	mandarin	76	5.8	4.0	0.81
8	1	apple	braeburn	178	7.1	7.8	0.92
9	1	apple	braeburn	172	7.4	7.0	0.89
10	1	apple	braeburn	166	6.9	7.3	0.93
11	1	apple	braeburn	172	7.1	7.6	0.92
12	1	apple	braeburn	154	7.0	7.1	0.88
13	1	apple	golden_delicious	164	7.3	7.7	0.70
14	1	apple	golden_delicious	152	7.6	7.3	0.69
15	1	apple	golden_delicious	156	7.7	7.1	0.69
16	1	apple	golden_delicious	156	7.6	7.5	0.67
17	1	apple	golden_delicious	168	7.5	7.6	0.73
18	1	apple	cripps_pink	162	7.5	7.1	0.83
19	1	apple	cripps_pink	162	7.4	7.2	0.85
20	1	apple	cripps_pink	160	7.5	7.5	0.86

These four columns contain the features of each data instance (sample)

The scale for the (simplistic) `color_score` feature used in the fruit dataset



Creating Training and Testing Sets





Applied Machine Learning

Examining the Data

Kevyn Collins-Thompson

Associate Professor of Information & Computer Science
University of Michigan


Some reasons why looking at the data initially is important

Examples of incorrect or missing feature values

- **Inspecting feature values may help identify what cleaning or preprocessing still needs to be done once you can see the range or distribution of values that is typical for each attribute.**
- **You might notice missing or noisy data, or inconsistencies such as the wrong data type being used for a column, incorrect units of measurements for a particular column, or that there aren't enough examples of a particular class.**
- **You may realize that your problem is actually solvable without machine learning.**


	fruit_label	fruit_name	fruit_subtype	mass	width	height	color_score
0	1	apple	granny_smith	192	8.4	7.3	0.55
1	1	apple	granny_smith	180	8.0	6.8	0.59
2	1	apple	granny_smith	176	7.4	7.2	192
3	2	mandarin	mandarin	86	6.2	4.7	0.80
4	2	mandarin	mandarin	84	6.0	4.6	0.79
5	2	mandarin	apple	80	5.8	4.3	0.77
6	2	mandarin	mandarin	80	5.9	4.3	0.81
7	2	mandarin	mandarin	76	5.8	4.0	0.81
8	1	apple	braeburn	178	7.1	7.8	0.92
9	1	apple	braeburn		7.4	7.0	0.89
10	1	apple	braeburn		6.9	7.3	0.93
11	1	apple	braeburn		7.1	7.6	0.92
12	1	apple	braeburn		7.0	7.1	0.88
13	1	apple	golden_delicious	164	7.3	7.7	0.70
14	1	apple	golden_delicious	152	7.6	7.3	0.69

A pairwise feature scatterplot visualizes the data using all possible pairs of features, with one scatterplot per feature pair, and histograms for each feature along the diagonal.



Individual scatterplot plotting all fruits by their **height** and **color_score**.

Colors represent different fruit classes.



In [27]: fruits

Out[27]:

	fruit_label	fruit_name	fruit_subtype	mass	width	height	color_score
0	1	apple	granny_smith	192	8.4	7.3	0.55
1	1	apple	granny_smith	180	8.0	6.8	0.59
2	1	apple	granny_smith	176	7.4	7.2	0.60
3	2	mandarin	mandarin	86	6.2	4.7	0.80
4	2	mandarin	mandarin	84	6.0	4.6	0.79
5	2	mandarin	mandarin	80	5.8	4.3	0.77
6	2	mandarin	mandarin	80	5.9	4.3	0.81
7	2	mandarin	mandarin	76	5.8	4.0	0.81
8	1	apple	braeburn	178	7.1	7.8	0.92
9	1	apple	braeburn	172	7.4	7.0	0.89
10	1	apple	braeburn	166	6.9	7.3	0.93
11	1	apple	braeburn	172	7.1	7.6	0.92
12	1	apple	braeburn	154	7.0	7.1	0.88
13	1	apple	golden_delicious	164	7.3	7.7	0.70
14	1	apple	golden_delicious	152	7.6	7.3	0.69
15	1	apple	golden_delicious	156	7.7	7.1	0.69
16	1	apple	golden_delicious	156	7.6	7.5	0.67

In [88]: fruits.shape

Out[88]: (59, 7)

In [88]: fruits.shape

Out[88]: (59, 7)

In [92]: X_train.shape

Out[92]: (44, 4)

In [93]: X_test.shape

Out[93]: (15, 4)

In [94]: y_train.shape

Out[94]: (44,)

In [95]: y_test.shape

Out[95]: (15,)

In [96]: X_train

Out[96]:

	height	width	mass	color_score
42	7.2	7.2	154	0.82
48	10.1	7.3	174	0.72
7	4.0	5.8	76	0.81
14	7.3	7.6	152	0.69
32	7.0	7.2	164	0.80
49	8.7	5.8	132	0.73
29	7.4	7.0	160	0.81
37	7.3	7.3	154	0.79
56	8.1	5.9	116	0.73
18	7.1	7.5	162	0.83
55	7.7	6.3	116	0.72
27	9.2	7.5	204	0.77
15	7.1	7.7	156	0.69
5	4.3	5.8	80	0.77
31	8.0	7.8	210	0.82
16	7.5	7.6	156	0.67

In [98]: y_train

Out[98]:

In [97]: X_test


Out[97]:

	height	width	mass	color_score
26	9.2	9.6	362	0.74
35	7.9	7.1	150	0.75
43	10.3	7.2	194	0.70
28	7.1	6.7	140	0.72
11	7.6	7.1	172	0.92
2	7.2	7.4	176	0.60
34	7.8	7.6	142	0.75
46	10.2	7.3	216	0.71
40	7.5	7.1	154	0.78
22	7.1	7.3	140	0.87
4	4.6	6.0	84	0.79
10	7.3	6.9	166	0.93
30	7.5	7.1	158	0.79
41	8.2	7.6	180	0.79
33	8.1	7.5	190	0.74
1	1			
12	1			
45	4			
24	3			
6	2			

In [99]: y_test


Out[99]:

26	3
35	3
43	4
28	3
11	1
2	1
34	3
46	4
40	3
22	1
4	2
10	1
30	3
41	3
33	3




```
from matplotlib import cm
cmap = cm.get_cmap('gnuplot')
scatter = pd.scatter_matrix(X_train, c= y_train, marker = 'o', s=40, hist_kwds={'bins':15}, figsize=(12,12), cmap=cmap)
```

A three-dimensional feature scatterplot



A three-dimensional feature scatterplot



```
from mpl_toolkits.mplot3d import Axes3D
fig = plt.figure()
ax = fig.add_subplot(111, projection = '3d')
ax.scatter(X_train['width'], X_train['height'], X_train['color_score'], c = y_train, marker = 'o', s=100)
ax.set_xlabel('width')
ax.set_ylabel('height')
ax.set_zlabel('color_score')
plt.show()
```



Applied Machine Learning

K-Nearest Neighbors Classification

Kevyn Collins-Thompson


Associate Professor of Information & Computer Science
University of Michigan

The k-Nearest Neighbor (k-NN) Classifier Algorithm

Given a training set X_{train} with labels y_{train} , and
given a new instance x_{test} to be classified:

1. Find the most similar instances (let's call them X_{NN}) to x_{test} that are in X_{train} .
2. Get the labels y_{NN} for the instances in X_{NN}
3. Predict the label for x_{test} by combining the labels y_{NN}
e.g. simple majority vote

A visual explanation of k-NN



Fruit dataset
Decision boundaries
with $k = 1$

A nearest neighbor algorithm needs four things specified

1. A distance metric
2. How many 'nearest' neighbors to look at?
3. Optional weighting function on the neighbor points
4. Method for aggregating the classes of neighbor points

A nearest neighbor algorithm needs four things specified

1. A distance metric

Typically Euclidean (Minkowski with $p = 2$)

2. How many 'nearest' neighbors to look at?

e.g. five

3. Optional weighting function on the neighbor points


Ignored

4. How to aggregate the classes of neighbor points


Simple majority vote

(Class with the most representatives among nearest neighbors)


K-nearest neighbors (k=1) for fruit dataset




K-nearest neighbors (k=5) for fruit dataset




K-nearest neighbors ($k=10$) for fruit dataset






K=1




K=5



K=10

How sensitive is k-NN classifier accuracy to the choice of 'k' parameter?



Fruit dataset
with 75%/25%
train-test split